

NLP__Assemment-Project_02

August 20, 2022

1 Project Name:- Help Twitter Combat Hate Speech Using NLP and Machine Learning

```
[1]: import warnings
warnings.filterwarnings('ignore')
```

```
[2]: # Import Require library

import pandas as pd
import re
from nltk.tokenize import TweetTokenizer
from nltk.corpus import stopwords
from string import punctuation
from collections import Counter
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import \
    accuracy_score, classification_report, recall_score, f1_score
```

```
[3]: # Set column width 150 for showing most data
pd.set_option('display.max_colwidth',150)
```

2 1. Load the tweets file using read_csv function from Pandas package.

```
[4]: data = pd.read_csv('TwitterHate.csv',usecols=['label','tweet'])
data.head(10)
```

```
[4]:   label \
0      0
1      0
2      0
3      0
4      0
5      0
```

```
6      0
7      0
8      0
9      0
```

```

                                tweet
0                                @user when a father is
dysfunctional and is so selfish he drags his kids into his dysfunction.  #run
1                                @user @user thanks for #lyft credit i can't use cause
they don't offer wheelchair vans in pdx.  #disapointed #getthanked
2
bihday your majesty
3                                #model    i love u
take with u all the time in urð ±!!! ð ð ð ð
ð |ð |ð |
4
factsguide: society now      #motivation
5                                [2/2] huge fan fare and big talking before they
leave. chaos and pay disputes when they get there. #allshowandnogo
6                                @user
camping tomorrow @user @user @user @user @user @user @user dannyâ |
7  the next school year is the year for exams.ð - can't think about that ð
#school #exams  #hate #imagine #actorslife #revolutionschool #girl
8                                we won!!! love the
land!!! #allin #cavs #champions #cleveland #clevelandcavaliers â |
9
@user @user welcome here !  i'm  it's so #gr8 !
```

```
[5]: # Check the shape of tha dataset
data.shape
```

```
[5]: (31962, 2)
```

```
[6]: # Check the information of dataset
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31962 entries, 0 to 31961
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  ---
0    label    31962 non-null    int64
1    tweet    31962 non-null    object
dtypes: int64(1), object(1)
memory usage: 499.5+ KB
```

```
[7]: # Check the missing value in dataset
data.isna().sum().any()
```

```
[7]: False
```

```
[8]: data['label'].value_counts()
```

```
[8]: 0    29720
     1    2242
     Name: label, dtype: int64
```

3 2. Get the tweets into a list for easy text cleanup and manipulation.

```
[9]: tweet = data['tweet'].values
```

```
[10]: tweet[:5]
```

```
[10]: array([' @user when a father is dysfunctional and is so selfish he drags his
kids into his dysfunction.    #run',
        '@user @user thanks for #lyft credit i can't use cause they don't offer
wheelchair vans in pdx.    #disappointed #getthanked",
        '  bihday your majesty',
        '#model  i love u take with u all the time in urð\x9f\x93±!!! ð\x9f\x98\
x99ð\x9f\x98\x8eð\x9f\x91\x84ð\x9f\x91\x85ð\x9f\x92!ð\x9f\x92!ð\x9f\x92!  ',
        ' factsguide: society now    #motivation'], dtype=object)
```

4 3. To cleanup:

4.1 i. Normalize the casing.

```
[11]: tweet_lower = [term.lower() for term in tweet]
```

```
[12]: tweet_lower[:5]
```

```
[12]: [' @user when a father is dysfunctional and is so selfish he drags his kids into
his dysfunction.    #run',
        '@user @user thanks for #lyft credit i can't use cause they don't offer
wheelchair vans in pdx.    #disappointed #getthanked",
        '  bihday your majesty',
        '#model  i love u take with u all the time in urð\x9f\x93±!!! ð\x9f\x98\x99ð\x
9f\x98\x8eð\x9f\x91\x84ð\x9f\x91\x85ð\x9f\x92!ð\x9f\x92!ð\x9f\x92!  ',
        ' factsguide: society now    #motivation']
```

4.2 ii. Using regular expressions, remove user handles. These begin with '@'.

```
[13]: tweet_nonuser = [re.sub('@\w+', '', term) for term in tweet_lower]
```

```
[14]: tweet_nonuser[:5]
```

```
[14]: [' when a father is dysfunctional and is so selfish he drags his kids into his  
dysfunction. #run',  
      " thanks for #lyft credit i can't use cause they don't offer wheelchair vans  
in pdx. #disappointed #getthanked",  
      ' bihday your majesty',  
      '#model i love u take with u all the time in urð\x9f\x93±!!! ð\x9f\x98\x99ð\x  
9f\x98\x8eð\x9f\x91\x84ð\x9f\x91\x85ð\x9f\x92!ð\x9f\x92!ð\x9f\x92| ',  
      ' factsguide: society now #motivation']
```

4.3 iii. Using regular expressions, remove URLs.

```
[15]: tweet_nonurl = [re.sub('\w+://\S+', '', term) for term in tweet_nonuser]
```

```
[16]: tweet_nonurl[:5]
```

```
[16]: [' when a father is dysfunctional and is so selfish he drags his kids into his  
dysfunction. #run',  
      " thanks for #lyft credit i can't use cause they don't offer wheelchair vans  
in pdx. #disappointed #getthanked",  
      ' bihday your majesty',  
      '#model i love u take with u all the time in urð\x9f\x93±!!! ð\x9f\x98\x99ð\x  
9f\x98\x8eð\x9f\x91\x84ð\x9f\x91\x85ð\x9f\x92!ð\x9f\x92!ð\x9f\x92| ',  
      ' factsguide: society now #motivation']
```

4.4 iv. Using TweetTokenizer from NLTK, tokenize the tweets into individual terms.

```
[17]: tweet_token = [TweetTokenizer().tokenize(sent) for sent in tweet_nonurl]
```

```
[18]: print(tweet_token[:5])
```

```
[['when', 'a', 'father', 'is', 'dysfunctional', 'and', 'is', 'so', 'selfish',  
'he', 'drags', 'his', 'kids', 'into', 'his', 'dysfunction', '.', '#run'],  
 ['thanks', 'for', '#lyft', 'credit', 'i', "can't", 'use', 'cause', 'they',  
"don't", 'offer', 'wheelchair', 'vans', 'in', 'pdx', '.', '#disappointed',  
'#getthanked'], ['bihday', 'your', 'majesty'], ['#model', 'i', 'love', 'u',  
'take', 'with', 'u', 'all', 'the', 'time', 'in', 'urð', '\x9f', '\x93', '±',  
'!', '!', '!', 'ð', '\x9f', '\x98', '\x99', 'ð', '\x9f', '\x98', '\x8e', 'ð',  
'\x9f', '\x91', '\x84', 'ð', '\x9f', '\x91', 'ð', '\x9f', '\x92', '|', 'ð',  
'\x9f', '\x92', '|', 'ð', '\x9f', '\x92', '|'], ['factsguide', ':', 'society',  
'now', '#motivation']]
```

4.5 v. Remove stop words.

```
[19]: stop_words = stopwords.words('english')
```

```
[20]: def remove_stopwords(input_word):  
        remove = [term for term in input_word if term not in stop_words]  
        return remove  
  
tweet_nonstopword = [remove_stopwords(sent) for sent in tweet_token]
```

```
[21]: print(tweet_nonstopword[:5])
```

```
[['father', 'dysfunctional', 'selfish', 'drags', 'kids', 'dysfunction', '.',  
 '#run'], ['thanks', '#lyft', 'credit', "can't", 'use', 'cause', 'offer',  
 'wheelchair', 'vans', 'pdx', '.', '#disapointed', '#getthanked'], ['bihday',  
 'majesty'], ['#model', 'love', 'u', 'take', 'u', 'time', 'urð', '\x9f', '\x93',  
 '±', '!', '!', '!', 'ð', '\x9f', '\x98', '\x99', 'ð', '\x9f', '\x98', '\x8e',  
 'ð', '\x9f', '\x91', '\x84', 'ð', '\x9f', '\x91', 'ð', '\x9f', '\x92', '||', 'ð',  
 '\x9f', '\x92', '||', 'ð', '\x9f', '\x92', '||'], ['factsguide', ':', 'society',  
 '#motivation']]
```

```
[22]: punct = list(punctuation)  
punct.extend(['...', '`', '"', "''", "..."])
```

```
[23]: def remove_punctuation(input_word):  
        remove = [term for term in input_word if term not in punct]  
        return remove  
  
tweet_nonpunctuation = [remove_punctuation(sent) for sent in tweet_nonstopword]
```

```
[24]: print(tweet_nonpunctuation[:5])
```

```
[['father', 'dysfunctional', 'selfish', 'drags', 'kids', 'dysfunction', '#run'],  
 ['thanks', '#lyft', 'credit', "can't", 'use', 'cause', 'offer', 'wheelchair',  
 'vans', 'pdx', '#disapointed', '#getthanked'], ['bihday', 'majesty'], ['#model',  
 'love', 'u', 'take', 'u', 'time', 'urð', '\x9f', '\x93', '±', 'ð', '\x9f',  
 '\x98', '\x99', 'ð', '\x9f', '\x98', '\x8e', 'ð', '\x9f', '\x91', '\x84', 'ð',  
 '\x9f', '\x91', 'ð', '\x9f', '\x92', '||', 'ð', '\x9f', '\x92', '||', 'ð', '\x9f',  
 '\x92', '||'], ['factsguide', 'society', '#motivation']]
```

4.6 vi. Remove redundant terms like ‘amp’, ‘rt’, etc.

```
[25]: context = ['amp', 'rt']
```

```
[26]: def stop_context(sent):  
        return [term for term in sent if term not in context]  
  
tweet_noncontext = [stop_context(tweet) for tweet in tweet_nonpunctuation]
```

```
[27]: print(tweet_noncontext[:5])
```

```
[['father', 'dysfunctional', 'selfish', 'drags', 'kids', 'dysfunction', '#run'],  
 ['thanks', '#lyft', 'credit', "can't", 'use', 'cause', 'offer', 'wheelchair',  
 'vans', 'pdx', '#disappointed', '#getthanked'], ['bihday', 'majesty'], ['#model',  
 'love', 'u', 'take', 'u', 'time', 'urð', '\x9f', '\x93', '±', 'ð', '\x9f',  
 '\x98', '\x99', 'ð', '\x9f', '\x98', '\x8e', 'ð', '\x9f', '\x91', '\x84', 'ð',  
 '\x9f', '\x91', 'ð', '\x9f', '\x92', '||', 'ð', '\x9f', '\x92', '||', 'ð', '\x9f',  
 '\x92', '||'], ['factsguide', 'society', '#motivation']]
```

4.7 vii. Remove ‘#’ symbols from the tweet while retaining the term.

```
[28]: def remove_symbol(sent):  
       return [re.sub('#', '', term) for term in sent]
```

```
[29]: tweet_nonsymbol = [remove_symbol(tweet) for tweet in tweet_noncontext]
```

```
[30]: print(tweet_nonsymbol[:5])
```

```
[['father', 'dysfunctional', 'selfish', 'drags', 'kids', 'dysfunction', 'run'],  
 ['thanks', 'lyft', 'credit', "can't", 'use', 'cause', 'offer', 'wheelchair',  
 'vans', 'pdx', 'disappointed', 'getthanked'], ['bihday', 'majesty'], ['model',  
 'love', 'u', 'take', 'u', 'time', 'urð', '\x9f', '\x93', '±', 'ð', '\x9f',  
 '\x98', '\x99', 'ð', '\x9f', '\x98', '\x8e', 'ð', '\x9f', '\x91', '\x84', 'ð',  
 '\x9f', '\x91', 'ð', '\x9f', '\x92', '||', 'ð', '\x9f', '\x92', '||', 'ð', '\x9f',  
 '\x92', '||'], ['factsguide', 'society', 'motivation']]
```

5 4. Extra cleanup by removing terms with a length of 1.

```
[31]: def clean_tweet(sent):  
       return [term for term in sent if len(term)>1]
```

```
[32]: tweet_clean = [clean_tweet(sent) for sent in tweet_nonsymbol]
```

```
[33]: print(tweet_clean[:5])
```

```
[['father', 'dysfunctional', 'selfish', 'drags', 'kids', 'dysfunction', 'run'],  
 ['thanks', 'lyft', 'credit', "can't", 'use', 'cause', 'offer', 'wheelchair',  
 'vans', 'pdx', 'disappointed', 'getthanked'], ['bihday', 'majesty'], ['model',  
 'love', 'take', 'time', 'urð'], ['factsguide', 'society', 'motivation']]
```

6 5. Check out the top terms in the tweets:

6.1 i. First, get all the tokenized terms into one large list.

```
[34]: tweet_list = []  
      for token in tweet_clean:  
          tweet_list.extend(token)
```

6.2 ii. Use the counter and find the 10 most common terms.

```
[35]: most_common_10 = Counter(tweet_list)  
      most_common_10.most_common(10)
```

```
[35]: [('love', 2748),  
        ('day', 2276),  
        ('happy', 1684),  
        ('time', 1131),  
        ('life', 1118),  
        ('like', 1047),  
        ("i'm", 1018),  
        ('today', 1013),  
        ('new', 994),  
        ('thankful', 946)]
```

7 6. Data formatting for predictive modeling:

7.1 i. Join the tokens back to form strings. This will be required for the vectorizers.

```
[36]: tweet_clean = [' '.join(term) for term in tweet_clean]
```

```
[37]: tweet_clean[:5]
```

```
[37]: ['father dysfunctional selfish drags kids dysfunction run',  
        "thanks lyft credit can't use cause offer wheelchair vans pdx disapointed  
        getthanked",  
        'bihday majesty',  
        'model love take time urð',  
        'factsguide society motivation']
```

7.2 ii. Assign x and y.

```
[38]: len(tweet_clean)
```

```
[38]: 31962
```

```
[39]: len(data['label'])
```

```
[39]: 31962
```

```
[40]: x = tweet_clean  
      y = data['label']
```

7.3 iii. Perform train_test_split using sklearn.

```
[41]: xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.2,random_state=42)
```

8 7. We'll use TF-IDF values for the terms as a feature to get into a vector space model.

8.1 i. Import TF-IDF vectorizer from sklearn.

```
[42]: from sklearn.feature_extraction.text import TfidfVectorizer
```

8.2 ii. Instantiate with a maximum of 5000 terms in your vocabulary.

```
[43]: vector = TfidfVectorizer(max_features=5000)
```

8.3 iii. Fit and apply on the train set.

```
[44]: xtrain_tfidf = vector.fit_transform(xtrain)
```

8.4 iv. Apply on the test set.

```
[45]: xtest_tfidf = vector.transform(xtest)
```

```
[46]: xtrain_tfidf.shape,xtest_tfidf.shape
```

```
[46]: ((25569, 5000), (6393, 5000))
```

9 8. Model building: Ordinary Logistic Regression

9.1 i. Instantiate Logistic Regression from sklearn with default parameters.

```
[47]: model = LogisticRegression()
```


9.2 ii. Fit into the train data.

```
[48]: model.fit(xtrain_tfidf,ytrain)
```

```
[48]: LogisticRegression()
```

9.3 iii. Make predictions for the train and the test set.

```
[49]: y_train_pred = model.predict(xtrain_tfidf)
      y_test_pred = model.predict(xtest_tfidf)
```

10 9. Model evaluation: Accuracy, recall, and f_1 score.

10.1 i. Report the accuracy on the train set.

```
[50]: print('Accuracy Score on Train Dataset:',accuracy_score(ytrain,y_train_pred))
```

Accuracy Score on Train Dataset: 0.9563142868317103

10.2 ii. Report the recall on the train set: decent, high, or low.

```
[51]: print('Recall Score on Train dataset:',
      ↪recall_score(ytrain,y_train_pred,average='weighted'))
```

Recall Score on Train dataset: 0.9563142868317103

10.3 iii. Get the f1 score on the train set.

```
[52]: print('F1 score on Train dataset:',
      ↪f1_score(ytrain,y_train_pred,average='weighted'))
```

F1 score on Train dataset: 0.9477746245284707

```
[53]: print('Classification Report on Train Dataset:')
      print(classification_report(ytest,y_test_pred))
```

Classification Report on Train Dataset:

	precision	recall	f1-score	support
0	0.95	1.00	0.97	5937
1	0.91	0.34	0.50	456
accuracy			0.95	6393
macro avg	0.93	0.67	0.73	6393
weighted avg	0.95	0.95	0.94	6393

11 10. Looks like you need to adjust the class imbalance, as the model seems to focus on the 0s.

11.1 i. Looks like you need to adjust the class imbalance, as the model seems to focus on the 0s.

```
[54]: model_imbalance = LogisticRegression(class_weight='balanced')
```

12 11. Train again with the adjustment and evaluate.

12.1 i. Train the model on the train set.

```
[55]: model_imbalance.fit(xtrain_tfidf,ytrain)
```

```
[55]: LogisticRegression(class_weight='balanced')
```

```
[56]: y_train_pred = model_imbalance.predict(xtrain_tfidf)
      y_test_pred = model_imbalance.predict(xtest_tfidf)
```

12.2 ii. Evaluate the predictions on the train set: accuracy, recall, and f₁ score.

```
[57]: print('Accuracy on Class Imbalance Train Dataset:
      ↪',accuracy_score(ytest,y_test_pred))
```

Accuracy on Class Imbalance Train Dataset: 0.923197246988894

```
[58]: print('Recall and F1 Score on Class Imbalance Train Dataset:')
      print(classification_report(ytest,y_test_pred))
```

```
Recall and F1 Score on Class Imbalance Train Dataset:
      precision    recall  f1-score   support

         0         0.98      0.93      0.96         5937
         1         0.48      0.80      0.60          456

   accuracy                   0.92         6393
  macro avg              0.73      0.87      0.78         6393
 weighted avg              0.95      0.92      0.93         6393
```

13 12. Regularization and Hyperparameter tuning:

13.1 i. Import GridSearch and StratifiedKFold because of class imbalance

```
[59]: from sklearn.model_selection import GridSearchCV, StratifiedKFold
```

13.2 ii. Provide the parameter grid to choose for 'C' and 'penalty' parameters.

```
[60]: params = {'C': [0.01, 0.1, 1, 10, 100],  
              'penalty': ['l1', 'l2']}
```

13.3 iii. Use a balanced class weight while instantiating the logistic regression.

```
[61]: model_gridsearch = LogisticRegression(class_weight='balanced')
```

14 13. Find the parameters with the best recall in cross validation.

14.1 i. Choose 'recall' as the metric for scoring.

14.2 ii. Choose stratified 4 fold cross validation scheme.

```
[62]: grid_search =  
      ↪ GridSearchCV(estimator=model_gridsearch, param_grid=params, cv=StratifiedKFold(4), n_jobs=-1, verbose=1,  
                    scoring='recall')
```

14.3 iii. Fit into the train set.

```
[63]: %%time  
      grid_search.fit(xtrain_tfidf, ytrain)
```

Fitting 4 folds for each of 10 candidates, totalling 40 fits
Wall time: 9.53 s

```
[63]: GridSearchCV(cv=StratifiedKFold(n_splits=4, random_state=None, shuffle=False),  
                  estimator=LogisticRegression(class_weight='balanced'), n_jobs=-1,  
                  param_grid={'C': [0.01, 0.1, 1, 10, 100], 'penalty': ['l1', 'l2']},  
                  scoring='recall', verbose=1)
```

15 14. What are the best parameters?

```
[64]: grid_search.best_estimator_
```

```
[64]: LogisticRegression(C=1, class_weight='balanced')
```

```
[65]: grid_search.best_params_
```

```
[65]: {'C': 1, 'penalty': 'l2'}
```

16 15. Predict and evaluate using the best estimator.

16.1 i. Use the best estimator from the grid search to make predictions on the test set.

```
[66]: y_test_grid_pred = grid_search.best_estimator_.predict(xtest_tfidf)
```

16.2 ii. What is the recall on the test set for the toxic comments?

```
[67]: print('Recall Score on Test Dataset:',  
        ↪recall_score(ytest,y_test_grid_pred,average='weighted'))
```

Recall Score on Test Dataset: 0.923197246988894

16.3 iii. What is the f₁ score?

```
[68]: print('F1 score:', f1_score(ytest,y_test_grid_pred,average='weighted'))
```

F1 score: 0.9318895024646877

```
[69]: print('Classification Report on GridSearchCV:')  
print(classification_report(ytest,y_test_grid_pred))
```

Classification Report on GridSearchCV:

	precision	recall	f1-score	support
0	0.98	0.93	0.96	5937
1	0.48	0.80	0.60	456
accuracy			0.92	6393
macro avg	0.73	0.87	0.78	6393
weighted avg	0.95	0.92	0.93	6393

```
[ ]:
```